ThreshSkip: Confidence-Thresholded Token Skipping for Video Diffusion Transformers

Synthetic prior work for the demo · public 2024-05-10

Abstract

We skip transformer blocks for spatio-temporal tokens of video diffusion models whose router confidence exceeds a fixed threshold, and copy those tokens forward. The threshold is constant across denoising steps.

3 Method

A router predicts a confidence for each token; tokens whose confidence exceeds a fixed threshold bypass the block. We evaluate on UCF-101 video generation and report 30% lower latency.